

CRWV

CoreWeave is the largest, fastest-scaling public pure-play AI cloud, a ~\$99.4B contracted backlog and privileged Nvidia acc...

Street PT **\$140.55 +92.0%**

CoreWeave

www.coreweave.com
Updated July 19, 2026 · 5 versions

CRWV PUBLIC DRAFT

- AI Cloud Infrastructure
- GPU Compute / GPU-as-a-Service
- Data Centers
- Cloud / Infra

33 ANALYSTS · CONSENSUS VIEW		
+92.0% upside to the Street target at 11.7x EV/Rev (LTM)		
PRICE	MKT CAP	STREET PT
\$73.21	\$39.9B	\$140.55
EV / REV	CONSENSUS	VS QQQ
11.7x	Buy	-65pp 1Y
All trading data as of the 2026-07-17 US market close (yfinance; EV/Rev recomputed from last close × shares + net debt). Figures covering a different window are labeled individually.		
FORWARD P/E	EV/REV 52-WK BAND	SHORT INTEREST
-50.0x	10.9x-18.7x	27.0%
consensus fwd EPS	price effect only; today's shares, net debt + LTM rev	% of float
NET DEBT	SBC % OF REV	
\$32.9B	9.6%	
cash – total debt	LTM, the dilution drag	

WHAT MATTERS	
KEY DEBATE	Is the ~\$99.4B backlog a durable annuity, or is a debt-funded, Microsoft/OpenAI-concentrated, GPU-depreciating model a fragile, circular AI-capex bet?
NEXT CATALYST	Q2 2026 results (guided to \$2.45-2.60B revenue) and the trajectory of 2026 capex (\$31-35B) and debt, the next read on whether growth still outruns burn.

BULL · BASE · BEAR

Bear Customer concentration: Microsoft was ~67% of FY2025 revenue and OpenAI/Meta dominate the backlog, the loss, renegotiation, or slowdown of any single anchor customer would be severe...	Base STREET \$140.55 (+92.0%) Street consensus Buy, 33 analysts; the \$140.55 target sits +92.0% above the \$73.21 close	Bull Customer diversification: OpenAI (~\$22B+), Meta (~\$35B), and Anthropic deals are pushing Microsoft toward less than 50% of future committed revenue...
---	--	---

BUSINESS OVERVIEW**IN ONE SENTENCE**

CoreWeave is an AI hyperscaler that rents out massive fleets of Nvidia GPUs to AI labs and tech giants, OpenAI, Microsoft, Meta, and Anthropic, on multi-year contracts. NASDAQ: CRWV; ~\$6.2B LTM revenue and a ~\$99.4B contracted backlog as of Q1 2026.

PLAIN ENGLISH

What it does, who pays

- CoreWeave runs giant data centers packed with Nvidia AI chips (GPUs).
- It rents that computing power to AI companies like OpenAI, Microsoft, and Meta.
- Customers sign multi-year contracts worth billions to reserve the capacity.
- It borrows heavily to build fast, betting contracted revenue pays it back.

THE TECHNICAL VERSION

How it actually works

- GPU-as-a-service: Nvidia GPUs in InfiniBand clusters tuned for AI training and inference.
- Revenue is mostly multi-year take-or-pay contracts, booked as a ~\$99.4B backlog (RPO).
- Capex is debt-financed, collateralized against signed customer contracts and the GPUs.
- High adjusted-EBITDA margins (~56%), but GAAP losses from interest and depreciation.

EXPLAINED SIMPLY

The analogy

- Imagine a landlord who builds warehouses full of AI supercomputers.
- Instead of buying their own, AI giants rent his machines on long leases.
- He takes on huge loans to build more, betting the rent more than covers it.

KEY TERMS, DECODED**GPU** *Graphics Processing Unit*

A specialized chip (here, Nvidia's) that does the massively parallel math AI models need to train and run.

Hyperscaler

A giant cloud provider operating at enormous scale (AWS, Microsoft Azure, Google Cloud); CoreWeave is an AI-focused challenger to them.

Neocloud

A new breed of cloud provider built specifically for AI/GPU workloads rather than general computing (e.g., CoreWeave, Nebius, IREN).

RPO *Remaining Performance Obligations*

The dollar value of signed contracts not yet recognized as revenue, i.e., the contracted backlog of future sales.

Take-or-pay

A contract where the customer must pay for reserved capacity whether or not they actually use it, giving the provider predictable revenue.

Capex *Capital Expenditures*

Money spent building long-lived assets, here, data centers, GPUs, and power infrastructure.

InfiniBand cluster

A high-speed network linking thousands of GPUs so they act as one machine, essential for training large AI models efficiently.

- (NASDAQ: CRWV) is an AI cloud-computing company headquartered in Livingston, New Jersey, that builds and operates GPU-dense data centers and rents that compute to AI developers, research labs, and enterprises.
- Founded in 2017 by commodities traders Michael Intrator, Brian Venturo, and Brannin McBee (with Peter Salanki as CTO), it began as crypto-mining venture 'Atlantic Crypto' before pivoting in 2018-19 to buy GPUs from distressed miners and rent out compute, rebranding as CoreWeave and reorienting around Nvidia chips for AI workloads.
- It is the largest and cleanest public pure-play 'neocloud,' differentiated by scale (roughly 30+ data centers and a fleet on the order of 250,000 GPUs), a privileged Nvidia relationship (Nvidia disclosed a ~7% stake in May 2025), and an asset-light model that leases data-center shells from operators like Equinix and Digital Realty.
- CoreWeave IPO'd in March 2025 at \$40/share and, after a sharp run, trades near \$96 as of the 2026-06-10 close (~\$42.8B market cap). The central debate is whether a ~\$99.4B contracted backlog is a durable annuity or a fragile, debt-funded, customer-concentrated bet on AI capex staying insatiable.

SOURCES: [COREWEAVE IR, IPO PRICING](#) · [COREWEAVE IR, Q1 2026 RESULTS](#) · [CNBC, NVIDIA ~7% STAKE / IPO CONTEXT](#)

EXECUTIVE LEADERSHIP

CO-FOUNDER, CHIEF EXECUTIVE OFFICER & CHAIRMAN

Michael Intrator

Former commodities trader; co-founded the company (as Atlantic Crypto) in 2017, led the 2018-19 pivot from crypto mining to AI cloud, and took CoreWeave public in March 2025.

CHIEF FINANCIAL OFFICER

Nitin Agrawal

CFO since March 2024; previously VP Finance of Google Cloud, CFO of Mapbox, and Finance Director of AWS's Compute Services division.

CO-FOUNDER & CHIEF STRATEGY OFFICER

Brian Venturo

Co-founder and former commodities trader; previously CTO, now Chief Strategy Officer leading capacity and platform strategy.

**CO-FOUNDER & CHIEF
DEVELOPMENT OFFICER****Brannin McBee**

Co-founder and former commodities trader; leads data-center development and capacity expansion.

**CO-FOUNDER & CHIEF TECHNOLOGY
OFFICER****Peter Salanki**

Co-founder and CTO; leads CoreWeave's engineering and infrastructure software stack.

PRODUCT & REVENUE MODEL

- CoreWeave sells GPU compute capacity ('GPU-as-a-service') primarily on large, multi-year committed contracts, frequently take-or-pay structures that underpin its enormous remaining-performance-obligation (RPO) backlog.
- It deploys Nvidia GPUs in InfiniBand-networked clusters tuned for AI training and inference, layered with its own orchestration and management software.
- The model is intensely capital-intensive and debt-financed: CoreWeave borrows against signed customer contracts (and the GPUs as collateral) to fund the buildout, then collects contracted revenue over the life of each deal.
- Revenue is highly concentrated, Microsoft was ~67% of FY2025 revenue, with OpenAI, Meta, Anthropic, Nvidia, IBM, and Jane Street among other anchor customers.
- Headline contracts include OpenAI (initial \$11.9B five-year deal in March 2025, since expanded toward ~\$22.4B+), Meta (~\$14.2B in September 2025, expanded by ~\$21B in April 2026 to ~\$35B through December 2032), and a new multi-year Anthropic agreement (April 2026). Growth is gated by power: CoreWeave surpassed 1 GW of active and >3.5 GW of total contracted power as of Q1 2026, targeting >1.7 GW active by year-end 2026.

SOURCES: [CNBC, OPENAI TO PAY COREWEAVE \\$11.9B OVER FIVE YEARS](#) ·

[COREWEAVE IR, EXPANDS OPENAI AGREEMENT BY UP TO \\$6.5B](#) · [COREWEAVE IR, Q1 2026 RESULTS \(BACKLOG / POWER\)](#)

RECENT NEWS

2026-07-15

Mizuho Cuts CoreWeave Price Target to \$100 from \$110 on Meta Compute Competition Risk

SOURCE: [CNBC](#)

Mizuho's PT cut formalizes analyst concern that Meta's compute-resale ambitions compress GPU cloud pricing for CoreWeave, even as Rosenblatt maintained a \$250 target citing no-sublease firewall language in the Meta contract. The wide analyst dispersion highlights the binary nature of the competitive threat debate.

2026-07-13

CoreWeave Stock Plunges ~35% as Meta Announces Plans to Resell Excess AI Compute Capacity

SOURCE: [REUTERS](#)

Meta's entry into AI compute resale (Meta Compute) introduces a direct competitive threat to CoreWeave's hyperscaler-adjacent cloud business, calling into question the durability of GPU cloud pricing and CoreWeave's ability to fill capacity beyond existing long-term contracts. The selloff reflects the market's binary concern about whether Meta's \$35.2B contract with CoreWeave includes firewall provisions against sublease competition.

2026-07-08	CoreWeave CTO Defends AI Cloud Moat Against Meta Competitive Threat SOURCE: BLOOMBERG CoreWeave co-founder/CTO Peter Salanki made the public case for purpose-built AI infrastructure (saturated networking, liquid cooling, predictive GPU failure) as a durable moat vs. CPU hyperscalers, management's rebuttal to the insourcing narrative.
2026-07-02	Meta Plans Competing AI Cloud Business, CoreWeave Bonds Slide SOURCE: BLOOMBERG Meta announced plans to sell excess AI compute to third-party customers, a direct competitive threat to CoreWeave's neocloud model. CoreWeave and Nebius shares each fell ~12%; credit markets questioned whether big-tech insourcing of GPU compute will commoditize neocloud economics.
2026-07-01	Meta Looks to Monetize Excess AI Compute, Overlapping CoreWeave's Customer Base SOURCE: TECHCRUNCH Meta, CoreWeave's customer with ~\$21B in committed spend, is now also a rival for enterprise AI workload buyers, compressing CoreWeave's TAM and raising concentration risk.
2026-06-18	CoreWeave Closes \$1.25B USD + €2B EUR Senior Notes, Part of \$8.5B Debt Build-Out SOURCE: SEC EDGAR The combined offering (9.625% USD / 8.500% EUR, both due 2032) demonstrates continued capital markets access at tighter spreads, funding CoreWeave's aggressive GPU data center expansion.
2026-06-12	CoreWeave to Join Nasdaq-100 Index on June 22 SOURCE: COREWEAVE INVESTOR RELATIONS Index inclusion ~15 months after the March 2025 IPO forces passive QQQ-tracking funds to buy CRWV shares, creating structural demand and cementing CoreWeave's large-cap AI infrastructure status.
2026-06-11	CoreWeave Prices \$3.5B Senior Unsecured Notes in First Euro-Denominated AI Junk Bond SOURCE: SEC EDGAR (8-K FILING) CoreWeave raised ~\$3.5B via a dual-currency notes offering (USD + EUR tranches, 6-year maturity), the first euro junk bond from a U.S. AI infrastructure company, to repay existing debt and fund operations, underscoring how aggressively the company is leveraging capital markets to fund GPU capacity expansion.

FINANCIALS & KEY METRICS

Reporting period: Q1 2026 (ended March 31, 2026); reported May 7, 2026. Each metric is tagged with the period it covers; LTM and as-of figures are labeled individually.

REVENUE

\$2.08B

▲ +111.6% YoY

Q1 2026

REVENUE BACKLOG RPO

\$99.4B

AS OF MARCH 31, 2026

ADJUSTED EBITDA

\$1.2B / 56% margin

Q1 2026

ADJUSTED OPERATING INCOME

\$21M / 1% margin

Q1 2026

GAAP NET LOSS

\$

740

Q1 2026

CAPITAL EXPENDITURES

\$7.70B

Q1 2026

CONTRACTED POWER

**>3.5 GW
contracted; >1
GW active**

AS OF Q1 2026

SEGMENT DETAIL, BOOKINGS & GUIDANCE


Revenue Guidance	Q2 2026 GUIDANCE	\$2.45-2.60B	Capex Guidance	FY2026 GUIDANCE	\$31-35B
Revenue Guidance	FY2026 GUIDANCE	\$12.0-13.0B			

- AI / Product:** CoreWeave called Q1 2026 its strongest bookings quarter ever, with revenue backlog reaching nearly \$100B (\$99.4B) on expanded deals with Meta (~\$21B, through Dec 2032), OpenAI, and a new multi-year Anthropic agreement to host development and inference of the Claude family of models; management said it now serves nine of the ten leading AI model providers. Revenue more than doubled to \$2.08B (+111.6% YoY) at a ~56% adjusted EBITDA margin, but adjusted operating margin compressed to just 1% as depreciation and interest scaled, and the GAAP net loss widened to \$740M.
- Demand:** Management framed demand as outstripping supply, with growth gated by power and capacity rather than orders ('capacity-constrained'). CFO Nitin Agrawal raised the low end of FY2026 capex guidance to \$31-35B, citing 'increases in component pricing.' The stock fell as much as ~10% after hours as Q2 revenue guidance (\$2.45-2.60B) came in below the ~\$2.69B consensus and the spending forecast rose, reviving the debate over cash burn and a debt load of nearly \$25B at quarter-end (with >\$20B of debt/equity raised year-to-date in 2026).


SOURCES: COREWEAVE IR, STRONG FIRST QUARTER 2026 RESULTS ·
SEC EDGAR, COREWEAVE Q1 2026 8-K EARNINGS PRESS RELEASE ·
CNBC, COREWEAVE Q1 2026: STOCK SINKS ON WEAK GUIDANCE, HIGHER CAPEX

REVENUE TRAJECTORY

ANNUAL REVENUE (\$B)



QUARTERLY REVENUE (\$B)



SOURCES: SEC EDGAR (as reported) · SEC EDGAR (as reported, FY2025 10-K) · yfinance / Yahoo Finance (live pull, 2026-06-10 close)

LAST 5 QUARTERS

	1Q25	2Q25	3Q25	4Q25	1Q26
REVENUE	\$982M	\$1.2B	\$1.4B	\$1.6B	\$2.1B
YOY					+111.6%
QOQ		+23.5%	+12.5%	+15.2%	+32.2%
GROSS MARGIN	73.3%	74.2%	73.0%	67.6%	65.5%

SOURCE: yfinance / Yahoo Finance (quarterly income statement) · calendar-quarter labels on fiscal period-ends; gross margin = gross profit / revenue. Most recent quarter (1Q26): 111.6% YoY revenue growth.

KEY RISKS & DEBATES

- !

Customer concentration: Microsoft was ~67% of FY2025 revenue and OpenAI/Meta dominate the backlog, the loss, renegotiation, or slowdown of any single anchor customer would be severe, even as management guides Microsoft toward less than half of future committed revenue.
- !

Debt load & capex intensity: CoreWeave closed Q1 2026 with nearly \$25B of debt and ~\$536M of quarterly interest expense, funding a \$31-35B 2026 capex plan, a debt-financed, contract-collateralized model highly exposed to rising rates, component-price inflation, and any execution slip.
- !

GPU depreciation & obsolescence: rapid Nvidia product cycles risk stranding capital in older GPUs; aggressive depreciation schedules pressure GAAP profitability, and useful-life assumptions are a live bear thesis.
- !

Circular AI financing: deep ties to Nvidia (~7% holder, chip supplier, and a customer) and capital flowing among Nvidia/OpenAI/CoreWeave fuel concerns that demand and funding are partly self-referential rather than fully independent end-demand.
- !

Profitability & cash burn: despite a ~56% adjusted EBITDA margin, CoreWeave is deeply GAAP loss-making (\$1.17B FY2025 net loss; \$740M in Q1 2026) with a thin 1% adjusted operating margin and heavy negative free cash flow as capex outruns operating cash.
- !

Power constraints & failed M&A: growth is gated by access to power and data-center capacity, the rejected \$9B Core Scientific acquisition removed a planned ~1.3 GW power source, forcing CoreWeave to secure capacity by other means amid intense competition from hyperscalers and neoclouds.


SOURCES: COREWEAVE IR, Q4 & FY2025 RESULTS (CONCENTRATION, NET LOSS) ·

CNBC, Q1 2026 WEAK GUIDANCE, HIGHER CAPEX · CNBC, CORE SCIENTIFIC SHAREHOLDERS REJECT \$9B DEAL

COMPARABLE COMPANIES

COMPANY	TICKER	TYPE	MKT CAP	EV / REV LTM	REV GROWTH YoY	GROSS MARGIN	FCF MARGIN LTM	RULE OF 40 growth + FCF
CoreWeave subject; largest public pure-play neocloud; ~\$99.4B backlog and privileged Nvidia access, offset by ~\$25B debt and Microsoft concentration	CRWV	Subject	\$39.9B	11.7x	111.6%	69.4%	-137.5%	-26
Nebius Group closest pure-play neocloud peer; far smaller, faster-growing off a tiny base, with a stronger, less-levered balance sheet after its 2025 raise	NBIS		\$45.1B	51.6x	683.9%	72.1%	-700.3%	-16
Applied Digital power-and-campus buildout sitting upstream of compute; a bet on converting contracted power into delivered GPU capacity, far smaller and less proven	APLD		\$7.4B	26.5x	139.3%	45.4%	-468.9%	-330
IREN bitcoin-miner-turned-neocloud, asset-heavy (owns its data centers); smaller and earlier-stage, mid-transition so reported growth screens flat	IREN		\$12.0B	18.2x	-0.0%	68.4%	-305.3%	-305
Oracle the most credible large-scale competitor; OCI is building AI capacity and signing mega-contracts with a fortress balance sheet and diversified software base CoreWeave can't match	ORCL		\$364.1B	7.4x	20.6%	65.8%	-36.4%	-16
Peer median	EX-SUBJECT			22.4x	80.0%	67.1%	-387.1%	-160

GROWTH VS MULTIPLE



SOURCES: Public-ticker multiples pulled live from [yfinance](#) / [Yahoo Finance](#), as of the 2026-07-17 market close; EV/Rev recomputed from last close × shares + net debt. Rev Growth (YoY), Gross Margin, and FCF Margin reflect each company's most recent reported quarter / LTM, so their as-of periods differ by company and from the price close above. Rule of 40 = revenue growth + FCF margin. **CoreWeave trades at 11.7x vs a 22.4x peer median, a 48% discount to the median.**

SWOT ANALYSIS

How CoreWeave stands out: CoreWeave is the largest, fastest-scaling public pure-play AI cloud, a ~\$99.4B contracted backlog and privileged Nvidia access set it apart from every neocloud peer, but it carries far more debt and customer concentration than the diversified hyperscalers, and trades at the lowest EV/Rev (12.2x as of the 2026-06-10 close) of its high-growth neocloud set.

STRENGTHS

INTERNAL EDGE

- **Privileged Nvidia relationship:** Nvidia is a ~7% shareholder, supplier, and customer; CoreWeave routinely gets early, large-scale access to the newest Nvidia GPUs in InfiniBand-clustered, AI-tuned deployments.
- **Enormous contracted backlog:** \$99.4B RPO as of Q1 2026 on multi-year take-or-pay deals with OpenAI, Meta, Microsoft, and Anthropic gives multi-year revenue visibility unmatched among neocloud peers.
- **Fastest-ever scaling at high cash margins:** Fastest cloud to \$5B annual revenue (FY2025 \$5.13B, +168%) at a ~56% adjusted EBITDA margin, real operating leverage on contracted capacity.

WEAKNESSES

INTERNAL GAP

- **Extreme customer concentration:** Microsoft was ~67% of FY2025 revenue and a handful of customers dominate the backlog, far more concentrated than any diversified cloud peer.
- **Heavy debt & negative GAAP economics:** ~\$25B of debt, ~\$536M quarterly interest expense, a \$740M Q1 2026 net loss, and only a 1% adjusted operating margin, profitability is consumed by interest and depreciation.
- **Capex / cash-burn dependence:** \$31-35B of 2026 capex funded largely by debt and equity leaves the model highly sensitive to financing conditions and Nvidia/component pricing.

OPPORTUNITIES

EXTERNAL UPSIDE

- **Customer diversification:** OpenAI (~\$22B+), Meta (~\$35B), and Anthropic deals are pushing Microsoft toward less than 50% of future committed revenue; CoreWeave now serves nine of ten leading AI providers.
- **Inference shift:** Newer contracts (e.g., Meta) are inference-weighted, a structurally larger, more durable, production-scale workload than one-off training runs.
- **Power-secured scale:** >3.5 GW of contracted power and a target of ~5 GW more by 2030 position it to capture continued AI compute demand if the buildout executes.

THREATS

EXTERNAL RISK

- **Hyperscaler & neocloud competition:** AWS/Azure/GCP and Oracle OCI self-fund capex and build custom silicon; Nebius, IREN, and Lambda compete for the same GPUs, customers, and power.
- **AI-spend / 'circular financing' reversal:** A pause in hyperscaler/lab AI capex, or scrutiny of Nvidia-linked circular funding, could hit both demand and CoreWeave's access to capital.
- **GPU obsolescence & rate risk:** Rapid Nvidia cycles threaten asset values, and a debt-heavy balance sheet is exposed to higher-for-longer rates and refinancing risk.

SOURCES: COREWEAVE IR, Q1 2026 RESULTS · YFINANCE / YAHOO FINANCE, LIVE COMPS (2026-06-10 CLOSE)

PRIMARY FILINGS (SEC EDGAR)

8-K	filed 2026-06-18 · period 2026-06-18	View on EDGAR >
8-K	filed 2026-06-11 · period 2026-06-11	View on EDGAR >
10-Q	filed 2026-05-08 · period 2026-03-31	View on EDGAR >
10-K	filed 2026-03-02 · period 2025-12-31	View on EDGAR >
10-Q	filed 2025-11-13 · period 2025-09-30	View on EDGAR >

SOURCE: SEC EDGAR · CIK 0001769628 · primary filings of record.

SLIDE-READY BULLETS

- 1 Largest public pure-play AI 'neocloud': Q1 2026 revenue more than doubled to \$2.08B (+111.6% YoY), with a contracted backlog (RPO) of \$99.4B, multi-year visibility unmatched among GPU-cloud peers.
- 2 Privileged Nvidia access is the moat: Nvidia is a ~7% holder, chip supplier, and customer, giving CoreWeave early, large-scale allocation of the newest GPUs in InfiniBand-clustered, AI-tuned deployments delivered faster than the generalist hyperscalers.
- 3 Customer base is diversifying off Microsoft (~67% of FY2025): OpenAI (~\$22B+), Meta (~\$35B through 2032), and a new Anthropic deal mean CoreWeave now serves nine of the ten leading AI model providers.
- 4 The bear case is the balance sheet: ~\$25B of debt, ~\$536M quarterly interest, a \$740M Q1 2026 GAAP net loss, and \$31-35B of 2026 capex make this a leveraged, capacity-constrained bet that AI demand stays insatiable.
- 5 Cheapest in its high-growth cohort: CoreWeave trades at ~12.2x EV/Revenue (as of the 2026-06-10 close) versus faster-growing neoclouds Nebius and Applied Digital, the discount the market assigns for concentration, leverage, and GPU-depreciation risk.

SMART DILIGENCE QUESTIONS

- Q1 How durable and renewable is the \$99.4B backlog? What share is true take-or-pay versus cancellable, what are the renewal economics as GPUs age, and how concentrated is it among OpenAI, Meta, and Microsoft, i.e., how much of the annuity survives a single anchor renegotiating?
- Q2 What is the unlevered, full-cycle return on a GPU cluster after interest and depreciation? With adjusted operating margin at just 1% and nearly \$25B of debt, at what GPU useful life and utilization does a contract actually generate positive free cash flow, and how sensitive is that to component-price inflation and rates?
- Q3 Where does the incremental power come from after the failed Core Scientific deal? What is the secured-vs-contracted-vs-speculative power pipeline to hit >1.7 GW active by year-end 2026 and ~5 GW by 2030, and what happens to the buildout and backlog conversion if power slips?

Sources: Company IR, earnings transcripts, Bloomberg, SEC filings, Sacra, Crunchbase, yfinance
 DRAFT for internal use only. Human review required before distribution.

ATLAS