

Databricks is the only at-scale, multi-cloud, open lakehouse that has turned its open-source roots into a >\$5.4B-run-rate, FCF-positive franchise growing >65%, uniquely positioned to ride enterprise AI by bringi...

Databricks

[www.databricks.com](https://www.databricks.com)

PRIVATE DRAFT

Updated July 19, 2026 · 4 versions

Data & AI Platform (Lakehouse)

Data Engineering & Warehousing

Machine Learning / GenAI

Cloud / Infra

|                                                              |                                                              |                                            |
|--------------------------------------------------------------|--------------------------------------------------------------|--------------------------------------------|
| <div>FREE CASH FLOW</div> <div>Positive</div> <div>LTM</div> | <div>VALUATION</div> <div>\$134B</div> <div>June 2026.</div> | <div>TOTAL RAISED</div> <div>\$22.0M</div> |
| <div>EMPLOYEES</div> <div>~14,500</div> <div>late 2025</div> |                                                              |                                            |

|               |                                                                                                                                                                                                                    |
|---------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WHAT MATTERS  |                                                                                                                                                                                                                    |
| KEY DEBATE    | Is a \$134B (soon possibly \$165-175B) valuation, roughly ~25-32x run-rate, justified by accelerating AI growth, or a rich late-cycle mark exposed to hyperscaler bundling and a Snowflake converging on its turf? |
| NEXT CATALYST | A new primary round reportedly in talks at \$165-175B (as early as mid-2026), with an eventual IPO Ghodsi now points toward ~2027 (he called 2026 'a terrible year' to list).                                      |

BUSINESS OVERVIEW

|                                                                                                                                                                                                                                                                                                         |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| IN ONE SENTENCE                                                                                                                                                                                                                                                                                         |
| Databricks is the private data-and-AI company behind the 'lakehouse', one platform unifying a company's data engineering, warehousing, analytics, and AI so teams build everything from dashboards to AI agents on a single copy of their data. >\$5.4B revenue run-rate (>65% YoY); \$134B last round. |

|                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| PLAIN ENGLISH                                                                                                                                                                                                                                                                                                                                                                                                        |
| What it does, who pays                                                                                                                                                                                                                                                                                                                                                                                               |
| <ul style="list-style-type: none"><li>Databricks gives big companies one place to store all their data and build AI on it.</li><li>It blends a cheap 'data lake' with a fast 'warehouse', the combo it calls a lakehouse.</li><li>Teams use it for dashboards, machine learning, and now AI chatbots and agents.</li><li>It's private, growing &gt;65% a year at a &gt;\$5.4B run-rate, valued at \$134B+.</li></ul> |

## THE TECHNICAL VERSION

How it actually works

- Lakehouse = open data-lake storage (Delta Lake/Iceberg) with warehouse-grade SQL and governance (Unity Catalog).
- Built on Apache Spark for distributed compute and MLflow for ML; runs on AWS, Azure, GCP.
- Mosaic AI handles fine-tuning, RAG, vector search, serving and agents (Agent Bricks); Claude is native.
- Consumption (DBU) pricing with >140% net retention; Lakebase adds serverless Postgres for agent workloads.

## EXPLAINED SIMPLY

The analogy

- Think of a giant, organized digital warehouse where a company keeps every piece of its data.
- Databricks is the warehouse plus the smart robots that find answers and build AI on those shelves.
- Instead of shipping your data out to an AI, you bring the AI into your own warehouse.

## KEY TERMS, DECODED

### Lakehouse

Databricks' architecture that combines a cheap data lake's open storage with a data warehouse's fast, governed querying, so you don't need two separate systems.

### Data warehouse

A structured, optimized database for fast business analytics and reporting on cleaned, organized data.

### Data lake

A low-cost store that holds raw data of any type (structured or not) at massive scale before it's organized.

### Apache Spark

The open-source engine (created by Databricks' founders) that processes huge datasets in parallel across many machines.

### Delta Lake

Databricks' open storage format that adds reliability (transactions, versioning) on top of a data lake.

### Unity Catalog

Databricks' governance layer that controls who can access which data and AI assets across the platform.

### ETL *Extract, Transform, Load*

Pulling data from sources, cleaning/reshaping it, and loading it where it can be analyzed.

### Run-rate (ARR) *Annualized Run-Rate / Annual Recurring Revenue*

A yearly-revenue estimate from annualizing current recurring or consumption revenue, useful for a fast-growing private company between reports.

### GenAI / LLM *Generative AI / Large Language Model*

AI models (like Claude) that generate text, code or answers, the foundation of the chatbots and agents Databricks helps companies build.

#### **NRR** *Net Revenue Retention*

How much more (or less) existing customers spend year over year; >100% means expansion beats churn, Databricks is >140%.

- Databricks, founded in 2013 by the UC Berkeley creators of Apache Spark (Ali Ghodsi, Matei Zaharia, Ion Stoica, Reynold Xin and others), is the company behind the 'Data Intelligence Platform', its branding of the lakehouse architecture that merges the cheap, open storage of a data lake with the structured query performance and governance of a data warehouse.
- Built on the open-source projects it created or stewards (Apache Spark, Delta Lake, MLflow, Unity Catalog) plus the open table format world it expanded into via Iceberg (the Tabular acquisition), the platform lets enterprises run data engineering, BI/SQL warehousing, streaming, and AI/ML on one governed copy of their data across AWS, Azure and Google Cloud.
- Since 2023 Databricks has pivoted hard into generative AI through its Mosaic AI stack (from the ~\$1.3B MosaicML acquisition), letting customers build, fine-tune and serve LLM-powered applications and agents directly on their data, and added a serverless operational database (Lakebase) via the ~\$1B Neon acquisition.
- It is the lead challenger to Snowflake in the data-platform wars and is headquartered in San Francisco.
- As of 2026 it remains private, with a revenue run-rate above \$5.4B (>65% YoY) and a \$134B post-money valuation, while reportedly in talks for a new round at \$165-175B.

SOURCES: **DATABRICKS NEWSROOM, \$5.4B RUN-RATE, >65% YOY (FEB 2026)** ·

**DATABRICKS NEWSROOM, SERIES L / \$4.8B RUN-RATE (DEC 2025)** · **SACRA, DATABRICKS PROFILE**

## EXECUTIVE LEADERSHIP

### CO-FOUNDER & CHIEF EXECUTIVE OFFICER

#### **Ali Ghodsi**

Co-creator of Apache Spark out of UC Berkeley; leads company vision and long-term strategy, and has publicly guided IPO timing toward ~2027.

### PRESIDENT, GLOBAL FIELD OPERATIONS

#### **Andy Kofoid**

Joined Databricks in August 2021 from Salesforce/DocuSign; owns the end-to-end customer journey from awareness to renewals.

### CHIEF FINANCIAL OFFICER

#### **Dave Conte**

CFO since October 2019 (previously CFO of Splunk); led the Series I-L mega-rounds and the free-cash-flow-positive milestone.

### CO-FOUNDER & CHIEF TECHNOLOGY OFFICER

#### **Matei Zaharia**

Original creator of Apache Spark and MLflow; UC Berkeley/Stanford; drives the technology and research roadmap.

### CO-FOUNDER & CHIEF ARCHITECT

#### **Reynold Xin**

Top Apache Spark contributor; architect of the platform's core engine.

### CO-FOUNDER & EXECUTIVE CHAIRMAN

#### **Ion Stoica**

Co-founder of Databricks and Anyscale; UC Berkeley professor leading the Sky Computing Lab; foundational to Spark's research origins.

## PRODUCT & REVENUE MODEL

---

- Databricks prices on consumption (usage), not seats: customers pay for the compute they consume, metered in Databricks Units (DBUs), on top of the underlying cloud, across AWS, Azure (first-party 'Azure Databricks') and Google Cloud.
- **The platform is layered:** (1) the lakehouse foundation, Delta Lake open storage plus Unity Catalog governance over a customer's data lake; (2) Databricks SQL / Data Warehousing, serverless SQL analytics that crossed a >\$1B run-rate in under four years; (3) Mosaic AI, model building, fine-tuning, RAG, vector search, model serving and Agent Bricks for production AI agents, including native access to Anthropic's Claude models via a five-year partnership; and (4) Lakebase, serverless Postgres (from Neon) as the operational/transactional database for AI agents, growing roughly 2x the pace of the warehousing business.
- Genie brings natural-language querying to enterprise data.
- Net revenue retention exceeds 140% via a land-and-expand motion; AI products alone now run above a \$1.4B annualized run-rate.

SOURCES: DATABRICKS & ANTHROPIC PARTNERSHIP (CLAUDE NATIVE) · DATABRICKS ACQUIRES NEON (LAKEBASE) · DATABRICKS NEWSROOM, \$5.4B RUN-RATE (GENIE / LAKEBASE)

---

## RECENT NEWS

---

2026-07-17

### Coatue Leads Databricks Funding Round at \$188 Billion Valuation

SOURCE: BLOOMBERG

Databricks raised ~\$3B in a new strategic round led by Coatue at a \$188B valuation, a 40% jump from its \$134B February 2026 Series L close, cementing its position as the leading independent AI/data platform. Funds are earmarked for acquisitions and expanding Unity AI Gateway (multi-AI governance), Genie (AI coworker for business data), and Lakebase (serverless Postgres for AI agents).

2026-07-17

### Databricks hits \$188B valuation, extending its run as AI's favorite second act

SOURCE: TECHCRUNCH

TechCrunch corroborates the Bloomberg report and adds context: this is Databricks' second fundraise of 2026 and the round has not yet closed (term sheet signed, close expected later in summer). The \$188B mark places Databricks among the most valuable private technology companies ever.

2026-06-25

### Databricks' Former AI Chief Naveen Rao Leaves to Launch Unconventional AI, Targeting 1,000x Power Efficiency

SOURCE: TECHCRUNCH

Naveen Rao, head of AI at Databricks, departed to found Unconventional AI (oscillator-based architecture for image generation). Signals leadership churn at a critical AI buildout moment and highlights talent-departure risk at late-stage private AI companies ahead of a potential Databricks IPO.

|            |                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2026-06-17 | <div><div>Databricks Doubles Down on Cybersecurity at Annual Conference, Unveils New Acquisition</div><div>SOURCE: AXIOS</div><div>Databricks announced a cybersecurity acquisition at its San Francisco conference, extending the Lakewatch SIEM platform launched in March, cybersecurity is becoming the company's third major product pillar alongside data/AI, expanding TAM and competitive surface vs. Splunk/Sentinel.</div></div> |
| 2026-06-16 | <div><div>Databricks Data Warehousing ARR Doubles to \$1.5B on AI Demand</div><div>SOURCE: BLOOMBERG</div><div>SQL/data warehousing ARR hit \$1.5B, doubling in a year, driven by AI workloads and Snowflake churn; the clearest public signal yet that Databricks is winning the data platform consolidation race.</div></div>                                                                                                            |
| 2026-06-16 | <div><div>Databricks Launches Unity AI Gateway to Prevent Enterprise AI Overspend</div><div>SOURCE: AXIOS</div><div>Unity AI Gateway adds spend limits, runaway-spend guardrails, and cross-provider model routing, positioning Databricks as the cost-control layer across the enterprise AI stack, not just the underlying data platform.</div></div>                                                                                    |
| 2026-06-08 | <div><div>Databricks in talks to raise at \$165-175B valuation, The Information reports</div><div>SOURCE: THE INFORMATION</div><div>A new round at this level would lift the valuation 23-31% above the \$134B Series L in under six months, underscoring sustained late-stage demand for AI-data infrastructure, and pushing the eventual IPO 'price to beat' higher.</div></div>                                                         |
| 2026-02-09 | <div><div>Databricks surpasses \$5.4B revenue run-rate, grows &gt;65% YoY; AI products at \$1.4B</div><div>SOURCE: DATABRICKS NEWSROOM</div><div>Acceleration (not deceleration) at this scale is rare, growth re-accelerated versus the &gt;55% reported in December 2025, and AI crossing a \$1.4B run-rate validates the GenAI pivot as a real revenue line.</div></div>                                                                |

FINANCIALS & KEY METRICS

Reporting period: Traction Snapshot, as of Feb 2026; reported As of 2026-06-11. Each metric is tagged with the period it covers; LTM and as-of figures are labeled individually.

|                                                                                                                                |                                                                        |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------|
| <div><div>REVENUE RUN RATE</div><div>&gt;\$5.4B annualized run-rate</div><div>&gt;65% YoY</div><div>AS OF FEB 2026</div></div> |                                                                        |
| <div><div>FREE CASH FLOW</div><div>Positive</div><div>LTM</div></div>                                                          | <div><div>FY 2025REVENUE</div><div>\$2.6B</div><div>FY2025</div></div> |

AI PRODUCTS

>\$1.4B annualized

AS OF FEB 2026

SQL WAREHOUSING

>\$1B annualized

crossed in under 4 years

RUN-RATE

NET RETENTION

>140%

AS OF FEB 2026

MILLION DOLLAR CUSTOMERS

>800 spending \$1M+; >70 spending \$10M+

AS OF FEB 2026

TOTAL CUSTOMERS

>20,000

orgs, incl. >60% of the Fortune 500

TRACTION SNAPSHOT — AS OF FEB 2026

LAST VALUATION

\$134B post-money

Series L

TRACTION SNAPSHOT — AS OF FEB 2026

SEGMENT DETAIL, BOOKINGS & GUIDANCE


Employees   TRACTION SNAPSHOT — AS OF FEB 2026   ~14,500

- AI / Product:** AI is now the fastest-growing part of the business: AI products crossed a >\$1.4B annualized run-rate by February 2026 (from >\$1B in December 2025). The Mosaic AI stack plus the five-year Anthropic partnership (Claude native in-platform), Agent Bricks for production agents, and Lakebase (serverless Postgres for agentic workloads, where >80% of databases were spun up by AI agents rather than humans) anchor the GenAI strategy, with Genie bringing natural-language querying to enterprise data.
- Demand:** A land-and-expand consumption model with >140% net retention drives durable growth; >800 customers now spend \$1M+ and >70 spend \$10M+ annually. The trajectory has accelerated through successive disclosures, \$2.6B in GAAP revenue for the year ended January 2025, then run-rate milestones of \$3.7B (July 2025), \$4.8B (December 2025) and >\$5.4B at >65% YoY (February 2026), while turning free-cash-flow positive over the trailing 12 months.

SOURCES: DATABRICKS NEWSROOM, \$5.4B RUN-RATE (FEB 2026) · DATABRICKS NEWSROOM, SERIES L / \$4.8B (DEC 2025) · CNBC, \$3.7B RUN-RATE, \$2.6B FY2025 REVENUE, >140% NRR

REVENUE TRAJECTORY

ANNUAL REVENUE (\$B)



SOURCES: Databricks (Series I press release) · CNBC / CFO Dave Conte · Databricks (Series L press release) · Databricks Newsroom

KEY RISKS & DEBATES

- ! Valuation vs. revenue (rich multiple): at a \$134B valuation (and reportedly \$165-175B in talks) on a >\$5.4B run-rate, Databricks trades at roughly ~25-32x run-rate revenue, far richer than public data peers, leaving little room for a growth stumble before an IPO.
- ! **Hyperscaler competition:** AWS (Redshift/SageMaker), Microsoft Fabric, and Google BigQuery bundle data + AI natively into the clouds Databricks runs on, and can subsidize or pre-install rival stacks, a structural channel and pricing threat.
- ! **Snowflake rivalry:** the direct lakehouse-vs-warehouse competitor is converging on Databricks' turf (open Iceberg tables, Cortex AI, its own LLM features), pressuring differentiation and pricing in the core analytics market.
- ! IPO timing & late-stage investor pressure: successive mega-rounds at ever-higher marks raise the bar the IPO must clear; Ghodsi points to ~2027, prolonging the liquidity overhang for employees and crossover backers.
- ! AI-platform execution & monetization durability: the AI run-rate (>\$1.4B) is growing fast but young, and consumption can be cut quickly in a downturn, proving Mosaic AI/Agent Bricks/Lakebase become durable, high-margin lines (not experimental spend) is unproven at scale.
- ! **Open-source commoditization:** Databricks' moat is built on open projects it no longer solely controls (Spark, Delta Lake, and now Iceberg via Tabular), open table formats and commoditized compute could let customers replicate core capabilities or migrate workloads more cheaply over time.

SOURCES: TECHCRUNCH, SERIES L AT \$134B · THE INFORMATION, DATABRICKS IN TALKS AT \$165-175B · SACRA, DATABRICKS PROFILE (COMPETITION / MODEL)

COMPARABLE COMPANIES

| COMPANY                                                                                                                                                                                                                                                                                                                                                       | TICKER | TYPE | MKT CAP | EV /<br>REV<br>LTM | REV<br>GROWTH<br>YoY | GROSS<br>MARGIN | FCF<br>MARGIN<br>LTM | RULE<br>OF 40<br>growth +<br>FCF |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------|------|---------|--------------------|----------------------|-----------------|----------------------|----------------------------------|
| <div><div>Databricks</div><div>subject (PRIVATE); \$134B last round (Series L, Feb 2026), in talks at \$165-175B, roughly ~25-32x its &gt;\$5.4B run-rate, a premium to every public data peer</div><div>Subject</div><div>&gt;65%</div></div>                                                                                                                |        |      |         |                    |                      |                 |                      |                                  |
| <div><div>Palantir</div><div>an AI-platform comp on narrative and multiple, not architecture; its rich multiple frames how the public market prices enterprise-AI hyper-growth, a reference point for where Databricks could list</div><div>PLTR</div><div>\$304.0B</div><div>56.7x</div><div>84.7%</div><div>84.1%</div><div>33.6%</div><div>118</div></div> |        |      |         |                    |                      |                 |                      |                                  |
| <div><div>Datadog</div><div>a consumption-pricing, land-and-expand SaaS comp (observability) often used to benchmark Databricks' usage-based model and &gt;140% net retention</div><div>DDOG</div><div>\$92.1B</div><div>24.1x</div><div>32.2%</div><div>79.9%</div><div>25.5%</div><div>58</div></div>                                                       |        |      |         |                    |                      |                 |                      |                                  |
| <div><div>Snowflake</div><div>the direct lakehouse-vs-warehouse rival; public and growing ~30%, while private Databricks grows roughly 2x faster, both now converging on each other's turf (Iceberg, native AI)</div><div>SNOW</div><div>\$93.2B</div><div>18.5x</div><div>33.5%</div><div>67.1%</div><div>34.6%</div><div>68</div></div>                     |        |      |         |                    |                      |                 |                      |                                  |
| <div><div>MongoDB</div><div>operational/document-database leader rather than an analytics platform, but increasingly overlapping with Databricks' new Lakebase (serverless Postgres for AI apps)</div><div>MDB</div><div>\$25.1B</div><div>8.7x</div><div>25.2%</div><div>72.0%</div><div>19.9%</div><div>45</div></div>                                      |        |      |         |                    |                      |                 |                      |                                  |
| Peer median EX-SUBJECT                                                                                                                                                                                                                                                                                                                                        |        |      |         | 21.3x              | 32.9%                | 76.0%           | 29.6%                | 63                               |

GROWTH VS MULTIPLE





**SOURCES:** Public-ticker multiples pulled live from [yfinance / Yahoo Finance](#), as of the 2026-07-17 market close; EV/Rev recomputed from last close × shares + net debt. Rev Growth (YoY), Gross Margin, and FCF Margin reflect each company's most recent reported quarter / LTM, so their as-of periods differ by company and from the price close above. Rule of 40 = revenue growth + FCF margin.

## SWOT ANALYSIS

**How Databricks stands out:** Databricks is the only at-scale, multi-cloud, open lakehouse that has turned its open-source roots (Spark, Delta, Iceberg) into a >\$5.4B-run-rate, FCF-positive franchise growing >65%, uniquely positioned to ride enterprise AI by bringing models to a customer's governed data rather than the reverse. Public proxy Snowflake trades at ~16.5x EV/Rev (2026-06-10 close); Databricks' last private round implies a richer ~25-32x run-rate mark.

### STRENGTHS

INTERNAL EDGE

- **Open lakehouse + multi-cloud reach:** Runs natively on AWS, Azure (first-party) and GCP on open formats (Delta Lake, Iceberg), avoiding lock-in and spanning a customer's entire data + AI estate on one governed copy.
- **Hypergrowth at scale with FCF:** >\$5.4B run-rate growing >65% YoY while FCF-positive over the trailing 12 months, a rare combination of scale, acceleration and profitability.
- **AI distribution & data gravity:** Mosaic AI + native Claude (Anthropic) + Agent Bricks + Lakebase let 20,000+ enterprises build AI on data they already store in Databricks; >140% net retention shows the engine compounding.

### WEAKNESSES

INTERNAL GAP

- **Rich private multiple:** ~25-32x run-rate revenue at \$134B (more at \$165-175B) leaves little margin for error and a demanding bar for the eventual IPO.
- **Consumption-model cyclicality:** Usage-based revenue can be throttled quickly by customers in a downturn, making the run-rate more volatile than seat-based SaaS.
- **Liquidity overhang:** No public stock and a delayed IPO (2027+) means employees and crossover investors face prolonged illiquidity, managed via periodic tender offers.

### OPPORTUNITIES

EXTERNAL UPSIDE

- **Operational-data TAM (Lakebase):** The Neon-powered serverless Postgres extends Databricks from analytics into live application/agent data, a large new market growing ~2x the warehousing business.
- **Enterprise agentic AI:** Agent Bricks + Genie position Databricks to be the platform of record for production AI agents over enterprise data as that workload explodes.
- **IPO optionality:** A marquee crossover syndicate and 'IPO-ready' governance set up a high-profile listing (potentially 2027) that could re-rate the whole category.


THREATS

EXTERNAL RISK

- **Hyperscaler bundling:** AWS, Microsoft Fabric and Google BigQuery bundle competing data+AI stacks into the clouds Databricks depends on, and can compete on price and distribution.
- **Snowflake convergence:** Snowflake is adopting open tables (Iceberg) and native AI (Cortex), eroding Databricks' openness and AI differentiation in the core market.
- **Open-source commoditization:** The open formats that are Databricks' moat (Spark, Delta, Iceberg) also lower switching costs and let rivals replicate core capabilities.

SOURCES: DATABRICKS NEWSROOM, \$5.4B RUN-RATE ·  
YFINANCE / YAHOO FINANCE, LIVE PUBLIC COMPS (2026-06-10 CLOSE)

FUNDING HISTORY



Total raised: \$22.0M

NOTABLE INVESTORS



IPO READINESS

- Company describes itself as 'IPO-ready' (governance and financial structures in place), is free-cash-flow positive over the trailing 12 months, and runs a >\$5.4B revenue run-rate.
- CEO Ali Ghodsi has called 2026 'a terrible year' to go public (citing crowded mega-IPOs) and signals a likely listing as early as 2027.
- Continued private mega-rounds and employee tender offers provide interim liquidity.
- No registration statement filed publicly as of June 2026.

## SLIDE-READY BULLETS

- 1 The most consequential pre-IPO name in data infrastructure: >\$5.4B revenue run-rate growing >65% YoY (as of Feb 2026) while free-cash-flow positive, a rare scale-plus-acceleration-plus-profitability profile.
- 2 The lakehouse thesis: one open, multi-cloud platform (Delta Lake + Unity Catalog) unifies data engineering, SQL warehousing, and AI, so enterprises bring AI to their governed data instead of shipping data out.
- 3 AI is the fastest-growing line: AI products crossed a >\$1.4B annualized run-rate, powered by Mosaic AI, native Anthropic Claude models, Agent Bricks, and the Neon-based Lakebase for agentic workloads.
- 4 Valuation has compounded fast: \$62B (Series J, Dec 2024) → >\$100B (Series K, Aug 2025) → \$134B (Series L, Feb 2026), reportedly now in talks at \$165-175B, implying ~25-32x run-rate revenue.
- 5 Key debate: is that mark justified by accelerating AI growth and >140% net retention, or is it a rich late-cycle valuation exposed to hyperscaler bundling and a Snowflake converging on its turf, with an IPO Ghodsi points toward ~2027?

## SMART DILIGENCE QUESTIONS

- Q1 How durable and high-margin is the AI run-rate? Of the >\$1.4B AI annualized revenue, how much is production workloads (Mosaic AI serving, Agent Bricks, Lakebase) versus experimental spend that can be cut, and what is the gross-margin profile of AI consumption versus core warehousing as inference scales?
- Q2 What does an apples-to-apples public valuation look like? On GAAP revenue (not run-rate) and against Snowflake's ~16.5x EV/Rev (2026-06-10 close), what multiple does the \$134B-\$175B private mark imply, and what growth/margin trajectory must hold for an IPO to clear that bar without a down-round?
- Q3 How defensible is the lakehouse against bundling and convergence? As AWS/Microsoft Fabric/Google BigQuery bundle data+AI and Snowflake adopts Iceberg + native AI, where is Databricks actually winning and losing competitive deals, and how much do the open formats (Spark, Delta, Iceberg) help retention versus lower switching costs for customers to leave?

Sources: Company IR, earnings transcripts, Bloomberg, SEC filings, Sacra, Crunchbase, yfinance  
DRAFT for internal use only. Human review required before distribution.

**ATLAS**